

# Capstone Project

September 25, 2020

## 1 Capstone Project

### 1.1 Image classifier for the SVHN dataset

#### 1.1.1 Instructions

In this notebook, you will create a neural network that classifies real-world images digits. You will use concepts from throughout this course in building, training, testing, validating and saving your Tensorflow classifier model.

This project is peer-assessed. Within this notebook you will find instructions in each section for how to complete the project. Pay close attention to the instructions as the peer review will be carried out according to a grading rubric that checks key parts of the project instructions. Feel free to add extra cells into the notebook as required.

#### 1.1.2 How to submit

When you have completed the Capstone project notebook, you will submit a pdf of the notebook for peer review. First ensure that the notebook has been fully executed from beginning to end, and all of the cell outputs are visible. This is important, as the grading rubric depends on the reviewer being able to view the outputs of your notebook. Save the notebook as a pdf (File -> Download as -> PDF via LaTeX). You should then submit this pdf for review.

#### 1.1.3 Let's get started!

We'll start by running some imports, and loading the dataset. For this project you are free to make further imports throughout the notebook as you wish.

```
In [1]: import tensorflow as tf
        from scipy.io import loadmat
        import numpy as np
        import pandas as pd
        import matplotlib.pyplot as plt
        %matplotlib inline
        import seaborn as sns
        sns.set_style('white')
        from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping
        from tensorflow.keras.models import Sequential, load_model
        from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D, Dropout, Bat
```



For the capstone project, you will use the [SVHN dataset](#). This is an image dataset of over 600,000 digit images in all, and is a harder dataset than MNIST as the numbers appear in the context of natural scene images. SVHN is obtained from house numbers in Google Street View images.

- Y. Netzer, T. Wang, A. Coates, A. Bissacco, B. Wu and A. Y. Ng. “Reading Digits in Natural Images with Unsupervised Feature Learning”. NIPS Workshop on Deep Learning and Unsupervised Feature Learning, 2011.

Your goal is to develop an end-to-end workflow for building, training, validating, evaluating and saving a neural network that classifies a real-world image into one of ten classes.

In [2]: # Run this cell to load the dataset

```
train = loadmat('data/train_32x32.mat')
test = loadmat('data/test_32x32.mat')
```

Both `train` and `test` are dictionaries with keys `X` and `y` for the input images and labels respectively.

## 1.2 1. Inspect and preprocess the dataset

- Extract the training and testing images and labels separately from the train and test dictionaries loaded for you.
- Select a random sample of images and corresponding labels from the dataset (at least 10), and display them in a figure.
- Convert the training and test images to grayscale by taking the average across all colour channels for each pixel. *Hint: retain the channel dimension, which will now have size 1.*
- Select a random sample of the grayscale images and corresponding labels from the dataset (at least 10), and display them in a figure.

In [3]: # Load the training and testing data and labels from the dictionary

```
train_x, train_y, test_x, test_y = train['X'], train['y'], test['X'], test['y']
print(f'Training data shape: {train_x.shape}\n')
```

```
Training labels shape:{train_y.shape}\n\  
Test data shape:{test_x.shape}\n\  
Test labels shape:{test_y.shape}'
```

```
Training data shape: (32, 32, 3, 73257)  
Training labels shape:(73257, 1)  
Test data shape:(32, 32, 3, 26032)  
Test labels shape:(26032, 1)
```

```
In [4]: # Rearrange the train and test data to have number of examples 1st
```

```
train_x = np.moveaxis(train_x, -1, 0)  
test_x = np.moveaxis(test_x, -1, 0)  
  
# Print the shape of the data again  
print(f'Training data shape: {train_x.shape}\n\  
Training labels shape:{train_y.shape}\n\  
Test data shape:{test_x.shape}\n\  
Test labels shape:{test_y.shape}'
```

```
Training data shape: (73257, 32, 32, 3)  
Training labels shape:(73257, 1)  
Test data shape:(26032, 32, 32, 3)  
Test labels shape:(26032, 1)
```

```
In [5]: # Replace the class label 10 to 0
```

```
test_y[test_y==10] = 0  
train_y[train_y==10] = 0
```

```
In [6]: # Helper function to convert color images to gray and normalize it.
```

```
def color_to_normalized_gray(arr):  
    """Convert array of colored images to gray and normalize it's values  
    by dividing by 255.  
    baram: arr: numpy array of shape(num_images, height, width, 3)  
    return normailized array of gray images.  
    """  
    gray_images = np.mean(arr, 3, keepdims=True) / 255  
    return gray_images
```

```
In [7]: train_x_gray = color_to_normalized_gray(train_x)  
test_x_gray = color_to_normalized_gray(test_x)  
# Print the shape of the data again  
print(f'New training data shape: {train_x_gray.shape}\n\  
New est data shape:{test_x_gray.shape}\n')
```

```
New training data shape: (73257, 32, 32, 1)  
New est data shape:(26032, 32, 32, 1)
```

```

In [8]: def plot_images(images, nrows, ncols, true_labels, predictions=None):
        """ Helper function for plotting nrows * ncols images
            param: images: array of images to plot
            param: nrows: number of rows in the subplot
            param: ncols: number of columns in the subplot
            param: true_labels: true label values of the images
            param: predictions: predicted labels of the images
            return None
        """
        fig, axes = plt.subplots(nrows, ncols, figsize=(16, 2*nrows))
        for i, ax in enumerate(axes.flat):
            if images[i].shape == (32, 32, 3):
                ax.imshow(images[i])
            else:
                ax.imshow(images[i,:,:,:0], cmap="gray")


            if predictions is None:
                title = f"Label: {str(true_labels[i])}"
                ax.set_title(title, color='b')
            elif predictions[i] == true_labels[i]:
                title = f"Label: {str(true_labels[i])}, Pred: [{str(predictions[i])}]"
                ax.set_title(title, color='g')
            else:
                title = f"Label: {str(true_labels[i])}, Pred: [{str(predictions[i])}]"
                ax.set_title(title, color='r')

            ax.set_xticks([]); ax.set_yticks([])

In [9]: # Get random indices for color images
        np.random.seed(0)
        idx_train = np.random.choice(train_x.shape[0]-1, 10)
        print('Color Images Examples')
        plot_images(train_x[idx_train], 2, 5, train_y[idx_train])


```

Color Images Examples



```
In [10]: # get random indices for Gray preprocessed images
np.random.seed(0)
idx_test = np.random.choice(train_x_gray.shape[0]-1, 10)
print('Gray Images Examples')
plot_images(train_x_gray[idx_test], 2, 5, train_y[idx_test])
```

Gray Images Examples



### 1.3 2. MLP neural network classifier

- Build an MLP classifier model using the Sequential API. Your model should use only Flatten and Dense layers, with the final layer having a 10-way softmax output.
- You should design and build the model yourself. Feel free to experiment with different MLP architectures. *Hint: to achieve a reasonable accuracy you won't need to use more than 4 or 5 layers.*
- Print out the model summary (using the summary() method)
- Compile and train the model (we recommend a maximum of 30 epochs), making use of both training and validation sets during the training run.
- Your model should track at least one appropriate metric, and use at least two callbacks during training, one of which should be a ModelCheckpoint callback.
- As a guide, you should aim to achieve a final categorical cross entropy training loss of less than 1.0 (the validation loss might be higher).
- Plot the learning curves for loss vs epoch and accuracy vs epoch for both training and validation sets.
- Compute and display the loss and accuracy of the trained model on the test set.

```
In [11]: # Create MLP model
def get_mlp_model():
    mlp_model = Sequential([
        Flatten(input_shape=train_x_gray[0].shape),
        Dense(512, activation='relu'),
        Dense(256, activation='relu'),
        Dense(128, activation='relu'),
        Dense(10, activation='softmax')
    ])
    return mlp_model
```

```
In [12]: mlp_model = get_mlp_model()
```

```
In [13]: # Print model summary
mlp_model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 1024)	0
dense (Dense)	(None, 512)	524800
dense_1 (Dense)	(None, 256)	131328
dense_2 (Dense)	(None, 128)	32896
dense_3 (Dense)	(None, 10)	1290
Total params: 690,314		
Trainable params: 690,314		
Non-trainable params: 0		

```
In [14]: # Compile the model
mlp_model.compile(optimizer='adam',
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])
```

```
In [15]: # Create call backs
checkpoint = ModelCheckpoint(filepath='mlp_model.h5',
                             save_best_only=True,
                             save_weights_only=True,
                             monitor='val_loss',
                             mode='min',
                             verbose=1)

early_stop = EarlyStopping(patience=5,
                           monitor='loss',
                           mode='min')
```

```
In [16]: # Train the model
history = mlp_model.fit(train_x_gray,
                        train_y,
                        epochs=30,
                        validation_data=(test_x_gray, test_y),
                        batch_size=128,
                        callbacks=[checkpoint, early_stop])
```

Train on 73257 samples, validate on 26032 samples

Epoch 1/30

73088/73257 [=====>.] - ETA: 0s - loss: 1.9441 - accuracy: 0.3117

Epoch 00001: val\_loss improved from inf to 1.54935, saving model to mlp\_model.h5

73257/73257 [=====] - 45s 608us/sample - loss: 1.9426 - accuracy: 0.3117

Epoch 2/30

73088/73257 [=====>.] - ETA: 0s - loss: 1.2301 - accuracy: 0.60 - ETA: 0s

Epoch 00002: val\_loss improved from 1.54935 to 1.18767, saving model to mlp\_model.h5

73257/73257 [=====] - 43s 591us/sample - loss: 1.2292 - accuracy: 0.60

Epoch 3/30

73088/73257 [=====>.] - ETA: 0s - loss: 1.0423 - accuracy: 0.6748

Epoch 00003: val\_loss improved from 1.18767 to 1.06973, saving model to mlp\_model.h5

73257/73257 [=====] - 42s 577us/sample - loss: 1.0422 - accuracy: 0.6748

Epoch 4/30

73088/73257 [=====>.] - ETA: 0s - loss: 0.9467 - accuracy: 0.7053

Epoch 00004: val\_loss improved from 1.06973 to 1.01101, saving model to mlp\_model.h5

73257/73257 [=====] - 42s 579us/sample - loss: 0.9467 - accuracy: 0.7053

Epoch 5/30

73216/73257 [=====>.] - ETA: 0s - loss: 0.8780 - accuracy: 0.7276

Epoch 00005: val\_loss improved from 1.01101 to 1.00514, saving model to mlp\_model.h5

73257/73257 [=====] - 43s 584us/sample - loss: 0.8778 - accuracy: 0.7276

Epoch 6/30

73088/73257 [=====>.] - ETA: 0s - loss: 0.8209 - accuracy: 0.7443

Epoch 00006: val\_loss improved from 1.00514 to 0.95072, saving model to mlp\_model.h5

73257/73257 [=====] - 43s 583us/sample - loss: 0.8205 - accuracy: 0.7443

Epoch 7/30

73216/73257 [=====>.] - ETA: 0s - loss: 0.7957 - accuracy: 0.7516

Epoch 00007: val\_loss did not improve from 0.95072

73257/73257 [=====] - 42s 580us/sample - loss: 0.7960 - accuracy: 0.7516

Epoch 8/30

73088/73257 [=====>.] - ETA: 0s - loss: 0.7536 - accuracy: 0.7672

Epoch 00008: val\_loss improved from 0.95072 to 0.84198, saving model to mlp\_model.h5

73257/73257 [=====] - 43s 583us/sample - loss: 0.7536 - accuracy: 0.7672

Epoch 9/30

73088/73257 [=====>.] - ETA: 0s - loss: 0.7283 - accuracy: 0.7732

Epoch 00009: val\_loss did not improve from 0.84198

73257/73257 [=====] - 43s 594us/sample - loss: 0.7282 - accuracy: 0.7732

Epoch 10/30

73216/73257 [=====>.] - ETA: 0s - loss: 0.7011 - accuracy: 0.7806

Epoch 00010: val\_loss improved from 0.84198 to 0.80137, saving model to mlp\_model.h5

73257/73257 [=====] - 43s 583us/sample - loss: 0.7011 - accuracy: 0.7806

Epoch 11/30

73216/73257 [=====>.] - ETA: 0s - loss: 0.6746 - accuracy: 0.7904

Epoch 00011: val\_loss improved from 0.80137 to 0.77974, saving model to mlp\_model.h5

73257/73257 [=====] - 43s 591us/sample - loss: 0.6747 - accuracy: 0.7904

Epoch 12/30

73216/73257 [=====>.] - ETA: 0s - loss: 0.6607 - accuracy: 0.7941 ETA: 0s

Epoch 00012: val\_loss did not improve from 0.77974

```

73257/73257 [=====] - 43s 588us/sample - loss: 0.6607 - accuracy: 0.7
Epoch 13/30
73088/73257 [=====>.] - ETA: 0s - loss: 0.6357 - accuracy: 0.8009
Epoch 00013: val_loss improved from 0.77974 to 0.77085, saving model to mlp_model.h5
73257/73257 [=====] - 43s 586us/sample - loss: 0.6357 - accuracy: 0.8
Epoch 14/30
73216/73257 [=====>.] - ETA: 0s - loss: 0.6239 - accuracy: 0.8049
Epoch 00014: val_loss improved from 0.77085 to 0.76856, saving model to mlp_model.h5
73257/73257 [=====] - 43s 587us/sample - loss: 0.6238 - accuracy: 0.8
Epoch 15/30
73088/73257 [=====>.] - ETA: 0s - loss: 0.5983 - accuracy: 0.8133
Epoch 00015: val_loss improved from 0.76856 to 0.74679, saving model to mlp_model.h5
73257/73257 [=====] - 43s 586us/sample - loss: 0.5984 - accuracy: 0.8
Epoch 16/30
73088/73257 [=====>.] - ETA: 0s - loss: 0.5932 - accuracy: 0.8132
Epoch 00016: val_loss improved from 0.74679 to 0.73286, saving model to mlp_model.h5
73257/73257 [=====] - 43s 589us/sample - loss: 0.5932 - accuracy: 0.8
Epoch 17/30
73088/73257 [=====>.] - ETA: 0s - loss: 0.5783 - accuracy: 0.8185 ETA: 1
Epoch 00017: val_loss did not improve from 0.73286
73257/73257 [=====] - 43s 588us/sample - loss: 0.5784 - accuracy: 0.8
Epoch 18/30
73216/73257 [=====>.] - ETA: 0s - loss: 0.5626 - accuracy: 0.8248
Epoch 00018: val_loss did not improve from 0.73286
73257/73257 [=====] - 43s 590us/sample - loss: 0.5625 - accuracy: 0.8
Epoch 19/30
73216/73257 [=====>.] - ETA: 0s - loss: 0.5508 - accuracy: 0.8273 ETA: 1
Epoch 00019: val_loss improved from 0.73286 to 0.73007, saving model to mlp_model.h5
73257/73257 [=====] - 43s 584us/sample - loss: 0.5508 - accuracy: 0.8
Epoch 20/30
73088/73257 [=====>.] - ETA: 0s - loss: 0.5422 - accuracy: 0.8301
Epoch 00020: val_loss improved from 0.73007 to 0.71365, saving model to mlp_model.h5
73257/73257 [=====] - 43s 587us/sample - loss: 0.5425 - accuracy: 0.8
Epoch 21/30
73216/73257 [=====>.] - ETA: 0s - loss: 0.5259 - accuracy: 0.8333
Epoch 00021: val_loss improved from 0.71365 to 0.69172, saving model to mlp_model.h5
73257/73257 [=====] - 43s 587us/sample - loss: 0.5259 - accuracy: 0.8
Epoch 22/30
73216/73257 [=====>.] - ETA: 0s - loss: 0.5236 - accuracy: 0.8354
Epoch 00022: val_loss did not improve from 0.69172
73257/73257 [=====] - 42s 580us/sample - loss: 0.5236 - accuracy: 0.8
Epoch 23/30
73216/73257 [=====>.] - ETA: 0s - loss: 0.5099 - accuracy: 0.8381
Epoch 00023: val_loss did not improve from 0.69172
73257/73257 [=====] - 43s 583us/sample - loss: 0.5099 - accuracy: 0.8
Epoch 24/30
73216/73257 [=====>.] - ETA: 0s - loss: 0.5006 - accuracy: 0.8411 ETA: 0
Epoch 00024: val_loss did not improve from 0.69172

```



```

73257/73257 [=====] - 43s 580us/sample - loss: 0.5005 - accuracy: 0.8
Epoch 25/30
73216/73257 [=====>.] - ETA: 0s - loss: 0.4894 - accuracy: 0.8459
Epoch 00025: val_loss did not improve from 0.69172
73257/73257 [=====] - 43s 582us/sample - loss: 0.4893 - accuracy: 0.8
Epoch 26/30
73216/73257 [=====>.] - ETA: 0s - loss: 0.4831 - accuracy: 0.8471
Epoch 00026: val_loss did not improve from 0.69172
73257/73257 [=====] - 43s 586us/sample - loss: 0.4831 - accuracy: 0.8
Epoch 27/30
73216/73257 [=====>.] - ETA: 0s - loss: 0.4728 - accuracy: 0.8519
Epoch 00027: val_loss improved from 0.69172 to 0.69095, saving model to mlp_model.h5
73257/73257 [=====] - 42s 578us/sample - loss: 0.4727 - accuracy: 0.8
Epoch 28/30
73216/73257 [=====>.] - ETA: 0s - loss: 0.4678 - accuracy: 0.8523
Epoch 00028: val_loss did not improve from 0.69095
73257/73257 [=====] - 43s 584us/sample - loss: 0.4679 - accuracy: 0.8
Epoch 29/30
73216/73257 [=====>.] - ETA: 0s - loss: 0.4602 - accuracy: 0.8552
Epoch 00029: val_loss did not improve from 0.69095
73257/73257 [=====] - 43s 591us/sample - loss: 0.4601 - accuracy: 0.8
Epoch 30/30
73088/73257 [=====>.] - ETA: 0s - loss: 0.4507 - accuracy: 0.8567
Epoch 00030: val_loss did not improve from 0.69095
73257/73257 [=====] - 42s 579us/sample - loss: 0.4508 - accuracy: 0.8

```

```

In [17]: df = pd.DataFrame(history.history)
         df.tail(10)

```

```

Out[17]:
```


	loss	accuracy	val_loss	val_accuracy
20	0.525938	0.833354	0.691724	0.797941
21	0.523641	0.835374	0.736311	0.789183
22	0.509937	0.838091	0.707300	0.793331
23	0.500456	0.841135	0.702182	0.795943
24	0.489301	0.845912	0.698412	0.800323
25	0.483087	0.847073	0.731530	0.786993
26	0.472740	0.851878	0.690947	0.802359
27	0.467890	0.852274	0.707651	0.798748
28	0.460082	0.855222	0.696088	0.805624
29	0.450791	0.856642	0.707214	0.801206

```


In [18]: # Plot the loss for training and validation sets
         plt.plot(df.loss)
         plt.plot(df.val_loss)
         plt.title('Loss vs. epochs')
         plt.ylabel('Loss')
         plt.xlabel('Epoch')

```

```
plt.legend(['Training', 'Validation'], loc='upper right')
plt.show()
```



```
In [19]: # Plot the accuracy for training and validation sets
plt.plot(df.accuracy)
plt.plot(df.val_accuracy)
plt.title('Accuracy vs. epochs')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Training', 'Validation'], loc='lower right')
plt.show()
```



```
In [20]: test_loss, test_acc = mlp_model.evaluate(test_x_gray, test_y, verbose=0)
         print("Test loss: {:.3f}\nTest accuracy: {:.2f}%".format(test_loss, 100 * test_acc))
```

Test loss: 0.707

Test accuracy: 80.12%

```
In [21]: # Load the best model
         mlp_best = get_mlp_model()
         mlp_best.summary()
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
flatten_1 (Flatten)	(None, 1024)	0
dense_4 (Dense)	(None, 512)	524800
dense_5 (Dense)	(None, 256)	131328
dense_6 (Dense)	(None, 128)	32896
dense_7 (Dense)	(None, 10)	1290

```
=====
Total params: 690,314
Trainable params: 690,314
Non-trainable params: 0
-----
```

```
In [22]: # Compile the model
mlp_best.compile(optimizer='adam',
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])
mlp_best.load_weights('mlp_model.h5')

In [23]: test_loss, test_acc = mlp_best.evaluate(test_x_gray, test_y, verbose=0)
print("Test loss: {:.3f}\nTest accuracy: {:.2f}%".format(test_loss, 100 * test_acc))

Test loss: 0.691
Test accuracy: 80.24%
```

## 1.4 3. CNN neural network classifier

- Build a CNN classifier model using the Sequential API. Your model should use the Conv2D, MaxPool2D, BatchNormalization, Flatten, Dense and Dropout layers. The final layer should again have a 10-way softmax output.
- You should design and build the model yourself. Feel free to experiment with different CNN architectures. *Hint: to achieve a reasonable accuracy you won't need to use more than 2 or 3 convolutional layers and 2 fully connected layers.*
- The CNN model should use fewer trainable parameters than your MLP model.
- Compile and train the model (we recommend a maximum of 30 epochs), making use of both training and validation sets during the training run.
- Your model should track at least one appropriate metric, and use at least two callbacks during training, one of which should be a ModelCheckpoint callback.
- You should aim to beat the MLP model performance with fewer parameters!
- Plot the learning curves for loss vs epoch and accuracy vs epoch for both training and validation sets.
- Compute and display the loss and accuracy of the trained model on the test set.

```
In [24]: cnn_model = Sequential([
    Conv2D(16, 3, padding='SAME', activation='relu', input_shape=(
    MaxPooling2D(2),
    BatchNormalization(),
    Dropout(0.3),
    Conv2D(32, 3, padding='SAME', activation='relu'),
    MaxPooling2D(2),
    BatchNormalization(),
    Conv2D(64, 3, padding='SAME', activation='relu'),
    MaxPooling2D(2),
    Dropout(0.3),
```

```

        Flatten(),
        Dense(64, activation='relu'),
        Dense(10, activation='softmax')
    ])

    cnn_model.summary()

Model: "sequential_2"
-----
Layer (type)                 Output Shape              Param #
-----
conv2d (Conv2D)              (None, 32, 32, 16)       160
-----
max_pooling2d (MaxPooling2D) (None, 16, 16, 16)        0
-----
batch_normalization_2 (Batch Normalization) (None, 16, 16, 16)        64
-----
dropout_2 (Dropout)          (None, 16, 16, 16)        0
-----
conv2d_1 (Conv2D)            (None, 16, 16, 32)       4640
-----
max_pooling2d_1 (MaxPooling2D) (None, 8, 8, 32)          0
-----
batch_normalization_3 (Batch Normalization) (None, 8, 8, 32)        128
-----
conv2d_2 (Conv2D)            (None, 8, 8, 64)        18496
-----
max_pooling2d_2 (MaxPooling2D) (None, 4, 4, 64)          0
-----
dropout_3 (Dropout)          (None, 4, 4, 64)          0
-----
flatten_2 (Flatten)          (None, 1024)              0
-----
dense_10 (Dense)              (None, 64)                65600
-----
dense_11 (Dense)              (None, 10)                 650
=====
Total params: 89,738
Trainable params: 89,642
Non-trainable params: 96
-----

In [25]: cnn_model.compile(optimizer='adam',
                           loss='sparse_categorical_crossentropy',
                           metrics=['accuracy'])

In [26]: checkpoint_cnn = ModelCheckpoint(filepath='CNN.h5',
                                           save_best_only=True,

```

```

        save_weights_only=False,
        monitor='val_accuracy',
        mode='max',
        verbose=1)
early_stop_cnn = EarlyStopping(monitor='loss', patience=5, verbose=1)

In [27]: history_cnn = cnn_model.fit(train_x_gray,
                                     train_y,
                                     callbacks=[checkpoint_cnn, early_stop_cnn],
                                     batch_size=128,
                                     validation_data=(test_x_gray, test_y),
                                     epochs=30)

Train on 73257 samples, validate on 26032 samples
Epoch 1/30
73216/73257 [=====>.] - ETA: 0s - loss: 1.2097 - accuracy: 0.5945
Epoch 00001: val_accuracy improved from -inf to 0.76721, saving model to CNN.h5
73257/73257 [=====] - 367s 5ms/sample - loss: 1.2094 - accuracy: 0.5945
Epoch 2/30
73216/73257 [=====>.] - ETA: 0s - loss: 0.5900 - accuracy: 0.8188
Epoch 00002: val_accuracy improved from 0.76721 to 0.86497, saving model to CNN.h5
73257/73257 [=====] - 352s 5ms/sample - loss: 0.5899 - accuracy: 0.8188
Epoch 3/30
73216/73257 [=====>.] - ETA: 0s - loss: 0.4903 - accuracy: 0.8516
Epoch 00003: val_accuracy improved from 0.86497 to 0.87846, saving model to CNN.h5
73257/73257 [=====] - 349s 5ms/sample - loss: 0.4903 - accuracy: 0.8516
Epoch 4/30
73216/73257 [=====>.] - ETA: 0s - loss: 0.4396 - accuracy: 0.8665
Epoch 00004: val_accuracy improved from 0.87846 to 0.88591, saving model to CNN.h5
73257/73257 [=====] - 350s 5ms/sample - loss: 0.4397 - accuracy: 0.8665
Epoch 5/30
73216/73257 [=====>.] - ETA: 0s - loss: 0.4047 - accuracy: 0.8769
Epoch 00005: val_accuracy improved from 0.88591 to 0.88791, saving model to CNN.h5
73257/73257 [=====] - 352s 5ms/sample - loss: 0.4048 - accuracy: 0.8769
Epoch 6/30
73216/73257 [=====>.] - ETA: 0s - loss: 0.3834 - accuracy: 0.8840
Epoch 00006: val_accuracy improved from 0.88791 to 0.90005, saving model to CNN.h5
73257/73257 [=====] - 352s 5ms/sample - loss: 0.3834 - accuracy: 0.8840
Epoch 7/30
73216/73257 [=====>.] - ETA: 0s - loss: 0.3641 - accuracy: 0.8891
Epoch 00007: val_accuracy improved from 0.90005 to 0.90154, saving model to CNN.h5
73257/73257 [=====] - 350s 5ms/sample - loss: 0.3639 - accuracy: 0.8891
Epoch 8/30
73216/73257 [=====>.] - ETA: 0s - loss: 0.3508 - accuracy: 0.8934
Epoch 00008: val_accuracy did not improve from 0.90154
73257/73257 [=====] - 350s 5ms/sample - loss: 0.3508 - accuracy: 0.8934
Epoch 9/30
73216/73257 [=====>.] - ETA: 0s - loss: 0.3388 - accuracy: 0.8955

```

Epoch 00009: val\_accuracy improved from 0.90154 to 0.91065, saving model to CNN.h5  
73257/73257 [=====] - 346s 5ms/sample - loss: 0.3388 - accuracy: 0.8984  
Epoch 10/30  
73216/73257 [=====>.] - ETA: 0s - loss: 0.3283 - accuracy: 0.8984  
Epoch 00010: val\_accuracy did not improve from 0.91065  
73257/73257 [=====] - 354s 5ms/sample - loss: 0.3282 - accuracy: 0.8984  
Epoch 11/30  
73216/73257 [=====>.] - ETA: 0s - loss: 0.3184 - accuracy: 0.9039  
Epoch 00011: val\_accuracy did not improve from 0.91065  
73257/73257 [=====] - 348s 5ms/sample - loss: 0.3184 - accuracy: 0.9039  
Epoch 12/30  
73216/73257 [=====>.] - ETA: 0s - loss: 0.3134 - accuracy: 0.9054  
Epoch 00012: val\_accuracy did not improve from 0.91065  
73257/73257 [=====] - 352s 5ms/sample - loss: 0.3133 - accuracy: 0.9054  
Epoch 13/30  
73216/73257 [=====>.] - ETA: 0s - loss: 0.3023 - accuracy: 0.9082  
Epoch 00013: val\_accuracy improved from 0.91065 to 0.91276, saving model to CNN.h5  
73257/73257 [=====] - 350s 5ms/sample - loss: 0.3023 - accuracy: 0.9082  
Epoch 14/30  
73216/73257 [=====>.] - ETA: 0s - loss: 0.2967 - accuracy: 0.9102  
Epoch 00014: val\_accuracy improved from 0.91276 to 0.91457, saving model to CNN.h5  
73257/73257 [=====] - 348s 5ms/sample - loss: 0.2966 - accuracy: 0.9102  
Epoch 15/30  
73216/73257 [=====>.] - ETA: 0s - loss: 0.2868 - accuracy: 0.9129  
Epoch 00015: val\_accuracy did not improve from 0.91457  
73257/73257 [=====] - 352s 5ms/sample - loss: 0.2869 - accuracy: 0.9129  
Epoch 16/30  
73216/73257 [=====>.] - ETA: 0s - loss: 0.2823 - accuracy: 0.9134  
Epoch 00016: val\_accuracy improved from 0.91457 to 0.91726, saving model to CNN.h5  
73257/73257 [=====] - 354s 5ms/sample - loss: 0.2823 - accuracy: 0.9134  
Epoch 17/30  
73216/73257 [=====>.] - ETA: 0s - loss: 0.2774 - accuracy: 0.9142  
Epoch 00017: val\_accuracy improved from 0.91726 to 0.92356, saving model to CNN.h5  
73257/73257 [=====] - 354s 5ms/sample - loss: 0.2774 - accuracy: 0.9142  
Epoch 18/30  
73216/73257 [=====>.] - ETA: 0s - loss: 0.2681 - accuracy: 0.9179  
Epoch 00018: val\_accuracy did not improve from 0.92356  
73257/73257 [=====] - 358s 5ms/sample - loss: 0.2682 - accuracy: 0.9179  
Epoch 19/30  
73216/73257 [=====>.] - ETA: 0s - loss: 0.2676 - accuracy: 0.9170  
Epoch 00019: val\_accuracy did not improve from 0.92356  
73257/73257 [=====] - 359s 5ms/sample - loss: 0.2676 - accuracy: 0.9170  
Epoch 20/30  
73216/73257 [=====>.] - ETA: 0s - loss: 0.2607 - accuracy: 0.9194  
Epoch 00020: val\_accuracy did not improve from 0.92356  
73257/73257 [=====] - 355s 5ms/sample - loss: 0.2608 - accuracy: 0.9194  
Epoch 21/30  
73216/73257 [=====>.] - ETA: 0s - loss: 0.2553 - accuracy: 0.9221

```

Epoch 00021: val_accuracy did not improve from 0.92356
73257/73257 [=====] - 354s 5ms/sample - loss: 0.2553 - accuracy: 0.9227
Epoch 22/30
73216/73257 [=====>.] - ETA: 0s - loss: 0.2512 - accuracy: 0.9227
Epoch 00022: val_accuracy did not improve from 0.92356
73257/73257 [=====] - 359s 5ms/sample - loss: 0.2511 - accuracy: 0.9227
Epoch 23/30
73216/73257 [=====>.] - ETA: 0s - loss: 0.2481 - accuracy: 0.9236
Epoch 00023: val_accuracy did not improve from 0.92356
73257/73257 [=====] - 347s 5ms/sample - loss: 0.2481 - accuracy: 0.9236
Epoch 24/30
73216/73257 [=====>.] - ETA: 0s - loss: 0.2463 - accuracy: 0.9247
Epoch 00024: val_accuracy did not improve from 0.92356
73257/73257 [=====] - 348s 5ms/sample - loss: 0.2463 - accuracy: 0.9247
Epoch 25/30
73216/73257 [=====>.] - ETA: 0s - loss: 0.2413 - accuracy: 0.9248
Epoch 00025: val_accuracy did not improve from 0.92356
73257/73257 [=====] - 354s 5ms/sample - loss: 0.2412 - accuracy: 0.9248
Epoch 26/30
73216/73257 [=====>.] - ETA: 0s - loss: 0.2370 - accuracy: 0.9260
Epoch 00026: val_accuracy did not improve from 0.92356
73257/73257 [=====] - 354s 5ms/sample - loss: 0.2370 - accuracy: 0.9260
Epoch 27/30
73216/73257 [=====>.] - ETA: 0s - loss: 0.2357 - accuracy: 0.9264
Epoch 00027: val_accuracy did not improve from 0.92356
73257/73257 [=====] - 360s 5ms/sample - loss: 0.2357 - accuracy: 0.9264
Epoch 28/30
73216/73257 [=====>.] - ETA: 0s - loss: 0.2348 - accuracy: 0.9259
Epoch 00028: val_accuracy improved from 0.92356 to 0.92594, saving model to CNN.h5
73257/73257 [=====] - 353s 5ms/sample - loss: 0.2347 - accuracy: 0.9259
Epoch 29/30
73216/73257 [=====>.] - ETA: 0s - loss: 0.2291 - accuracy: 0.9279
Epoch 00029: val_accuracy did not improve from 0.92594
73257/73257 [=====] - 354s 5ms/sample - loss: 0.2291 - accuracy: 0.9279
Epoch 30/30
73216/73257 [=====>.] - ETA: 0s - loss: 0.2263 - accuracy: 0.9292
Epoch 00030: val_accuracy did not improve from 0.92594
73257/73257 [=====] - 352s 5ms/sample - loss: 0.2262 - accuracy: 0.9292

```

```

In [28]: df_cnn = pd.DataFrame(history_cnn.history)
         df_cnn.tail(10)

```

```


Out[28]:
```

	loss	accuracy	val_loss	val_accuracy
20	0.255262	0.922042	0.290167	0.919292
21	0.251078	0.922765	0.278119	0.923517
22	0.248052	0.923611	0.283485	0.921827
23	0.246316	0.924676	0.289625	0.919945




24	0.241195	0.924826	0.285024	0.922711
25	0.236968	0.926028	0.304888	0.914605
26	0.235690	0.926423	0.286659	0.920713
27	0.234709	0.925946	0.279802	0.925937
28	0.229093	0.927884	0.280062	0.925668
29	0.226215	0.929181	0.280152	0.925668

```
In [29]: # Plot the loss for training and validation sets
plt.plot(df_cnn.loss)
plt.plot(df_cnn.val_loss)
plt.title('Loss vs. epochs')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Training', 'Validation'], loc='upper right')
plt.show()
```



```
In [30]: # Plot the accuracy for training and validation sets
plt.plot(df_cnn.accuracy)
plt.plot(df_cnn.val_accuracy)
plt.title('Accuracy vs. epochs')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Training', 'Validation'], loc='lower right')
plt.show()
```



```
In [31]: test_loss, test_acc = cnn_model.evaluate(test_x_gray, test_y, verbose=0)
         print("Test loss: {:.3f}\nTest accuracy: {:.2f}%".format(test_loss, 100 * test_acc))
```

Test loss: 0.280

Test accuracy: 92.57%

```
In [25]: cnn_best = load_model('CNN.h5')
         cnn_best.summary()
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 16)	160
max_pooling2d (MaxPooling2D)	(None, 16, 16, 16)	0
batch_normalization_2 (Batch Normalization)	(None, 16, 16, 16)	64
dropout_2 (Dropout)	(None, 16, 16, 16)	0
conv2d_1 (Conv2D)	(None, 16, 16, 32)	4640

max_pooling2d_1 (MaxPooling2)	(None, 8, 8, 32)	0
-----		
batch_normalization_3 (Batch Normalization)	(None, 8, 8, 32)	128
-----		
conv2d_2 (Conv2D)	(None, 8, 8, 64)	18496
-----		
max_pooling2d_2 (MaxPooling2)	(None, 4, 4, 64)	0
-----		
dropout_3 (Dropout)	(None, 4, 4, 64)	0
-----		
flatten_2 (Flatten)	(None, 1024)	0
-----		
dense_10 (Dense)	(None, 64)	65600
-----		
dense_11 (Dense)	(None, 10)	650
=====		
Total params: 89,738		
Trainable params: 89,642		
Non-trainable params: 96		
-----		

```
In [26]: test_loss, test_acc = cnn_best.evaluate(test_x_gray, test_y, verbose=0)
        print("Test loss: {:.3f}\nTest accuracy: {:.2f}%".format(test_loss, 100 * test_acc))
```

Test loss: 0.280  
Test accuracy: 92.59%

## 1.5 4. Get model predictions

- Load the best weights for the MLP and CNN models that you saved during the training run.
- Randomly select 5 images and corresponding labels from the test set and display the images with their labels.
- Alongside the image and label, show each model's predictive distribution as a bar chart, and the final model prediction given by the label with maximum probability.

```
In [24]: # load MLP best model

        # Create a new model with the same architecture
        mlp_best = get_mlp_model()

        # Compile the model
        mlp_best.compile(optimizer='adam',
                        loss='sparse_categorical_crossentropy',
                        metrics=['accuracy'])

        # Load the model weights
        mlp_best.load_weights('mlp_model.h5')
```

```

In [27]: # Load CNN best model
         cnn_best = load_model('CNN.h5')

In [28]: # Get random indices for test
         np.random.seed(0)
         idx = np.random.choice(test_x_gray.shape[0]-1, 5)

         # Get test images and labels
         test_images = test_x_gray[idx]
         test_labels = test_y[idx]

         # Get MLP and CNN predictions for test images
         mlp_predictions = mlp_best.predict(test_images)
         cnn_predictions = cnn_best.predict(test_images)

In [29]: # Show the test images and the MLP model predictions
         print('MLP Test Examples')
         plot_images(test_images, 1, 5, test_labels, np.argmax(mlp_predictions, 1))

```

MLP Test Examples



```

In [30]: # Show the test images and the CNN model predictions
         print('CNN Test Examples')
         plot_images(test_images, 1, 5, test_labels, np.argmax(cnn_predictions, 1))

```

CNN Test Examples



```

In [31]: def plot_predictions(predictions, test_images, test_labels, title):
         """Function to plot the predictions and the categorical probabiltty disturibution.

```

```

    param: predictions: array of shape (num_images, categories)
    param: test_images: test images to plot
    param: test_labels: true labels
    param: title: the kind of model for plot tutle
    """
    # Show the images and a bar chart with the predictions probabiltly
    fig, axes = plt.subplots(5, 2, figsize=(16, 12))
    fig.subplots_adjust(hspace=0.4, wspace=-0.2)


    for i, (prediction, image, label) in enumerate(zip(predictions, test_images, test_labels)):
        predicted_label = np.argmax(prediction)
        image_title = f"Label: {label}, Pred: [{predicted_label}]"
        axes[i, 0].imshow(image[:, :, 0], cmap='gray')
        axes[i, 0].get_xaxis().set_visible(False)
        axes[i, 0].get_yaxis().set_visible(False)
        if label == predicted_label:
            axes[i, 0].set_title(image_title, color='g')
        else:
            axes[i, 0].set_title(image_title, color='r')
        axes[i, 1].bar(np.arange(10), prediction)
        axes[i, 1].set_xticks(np.arange(10))
        axes[i, 1].set_title(f"Categorical Distribution ({title} Model Predictions)")

    plt.show()


```

```
In [32]: plot_predictions(mlp_predictions, test_images, test_labels, 'MLP')
```


Label: [2], Pred: [3]




Label: [9], Pred: [9]




Label: [2], Pred: [3]



Label: [1], Pred: [1]





Label: [2], Pred: [2]




```
In [33]: plot_predictions(cnn_predictions, test_images, test_labels, 'CNN')
```


Label: [2], Pred: [2]





Label: [9], Pred: [9]



Label: [2], Pred: [2]



Label: [1], Pred: [1]



Label: [2], Pred: [2]

